

Multi-Model Cascade Router

Standalone printable deep-dive dossier

Multi-Model Cascade Router

Agent systems fail when all tasks are pushed through one model. The router treats models as resources with strengths, costs, quotas, cooldowns, and failure modes.

- This shows model governance in practice: models are resources with availability, cost, quality, quota, privacy, and failure characteristics.
- Cost-reduction claims should be treated as design goals unless fresh logs verify exact percentages.

What I had in mind

I did not want to treat model choice as taste. I wanted routing to be part of the architecture: cheap when possible, strong when necessary, local when privacy matters, and fallback-aware when a provider fails.

What I built

- ``ai_router.py``, ``cascade.py``, and ``litellm_config.yaml`` define model registries, routing strategies, tiering, quota tracking, fallback, race/cancel behavior, and budget-aware escalation.

Mechanism detail

- The router treats models as infrastructure resources with strengths, costs, quotas, cooldowns, privacy profiles and failure modes.
- Task strategy determines whether to prefer speed, quality, free models, local execution, code capability or paid escalation.
- The core safety value is graceful degradation: when one model or tier fails, the system has a policy for trying another path rather than silently returning weak work.

Architecture

- Task strategy selects speed, quality, free, code, creative, bulk, or local routing.
- Local models are tried before free APIs and paid tiers where appropriate.
- Quota/cooldown tracking blocks bad or unavailable routes.
- Quality gates decide whether to escalate.
- Free API models can be raced in parallel and slower calls cancelled after a sufficient answer.
- Outcomes are logged for later learning.

Evidence cluster

- AI router source with provider registry for Gemini, Groq, Cerebras, OpenRouter, DeepSeek, Together, Ollama/local models, and LiteLLM-backed calls.

- Cascade source implementing local/free/paid tiers, quota and cooldown tracking, quality gates, free-model racing, cancellation, paid-budget blocking, and outcome logging.
- LiteLLM configuration separating local T0, free API T1, and paid T2 tiers.
- Architecture notes describing routing principles, quality scoring, budget handling, and quota-aware fallbacks.

Claim-to-evidence map

The public version should be strong because it is bounded, not because it overclaims.

Claim	Evidence	Boundary
This shows model governance in practice: models are resources with availability, cost, quality, quota, privacy, and failure characteristics.	AI router source with provider registry for Gemini, Groq, Cerebras, OpenRouter, DeepSeek, Together, Ollama/local models, and LiteLLM-backed calls.; Cascade source implementing local/free/paid tiers, quota and cooldown tracking, quality gates, free-model racing, cancellation, paid-budget blocking, and outcome logging.	Cost-reduction claims should be treated as design goals unless fresh logs verify exact percentages.
The work is valuable because it shows mechanism, not surface polish.	Task strategy selects speed, quality, free, code, creative, bulk, or local routing.; Local models are tried before free APIs and paid tiers where appropriate.	Fresh public demos or benchmarks would be the next proof layer.

Senior-level signal

This shows model governance in practice: models are resources with availability, cost, quality, quota, privacy, and failure characteristics.

- Provider-neutral thinking: no single model is treated as magic.
- Governance thinking: cost, quality, privacy and quota are part of model selection.
- Measurement path: outcome logging creates a basis for later empirical routing decisions.

Skeptical reviewer questions

- What is actually built here, and what is still design? Answer: the status and boundary are stated directly inside the Multi-Model Cascade Router case.
- Which private evidence is intentionally not public? Answer: local paths, credentials, private channel details and confidential raw material are excluded.
- Why is this relevant? Answer: the case shows a reusable component for agents, tool use, memory, routing, validation or high-stakes governance.
- What would be the next stronger proof? Answer: synthetic demo, audit trace, benchmark or external review, depending on the case.

Lessons and boundaries

Cost-reduction claims should be treated as design goals unless fresh logs verify exact percentages.

- A strong agent stack should degrade gracefully when models fail.

- Quality gates should decide escalation, not model status alone.
- Cost and privacy are part of capability routing.
- Outcome logging is the seed of empirical model governance.

Next hardening / next project step

- Run a fresh benchmark across representative task classes.
- Publish aggregate cost/latency/quality results without private provider details.
- Add policy labels for private, public, cheap, fast, and high-stakes tasks.
- Use the router as the basis for a secure-agent evaluation project.
- Run a benchmark on synthetic tasks: code repair, research summary, planning, long-context extraction and low-cost bulk routing.
- Publish aggregate latency/cost/quality results without exposing private provider details.
- Add policy labels so high-stakes or private tasks can be forced into safer routes.